

In [104..

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from optbinning import OptimalBinning
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from IPython.display import Image
import pydotplus
df = pd.read_csv('2022.csv')
```

# Logistic Regression

In [64]:

```
df = df.rename(columns = {"Happiness score":"Score"})
df['Is.Happy'] = [1 if each>np.mean(df.Score) else 0 for each in df.Score]

feature_cols = ['Explained by: GDP per capita', 'Explained by: Social support', 'Explained by: Healthy life exp
X = df[feature_cols].values
y = df['Is.Happy']
```

In [65]:

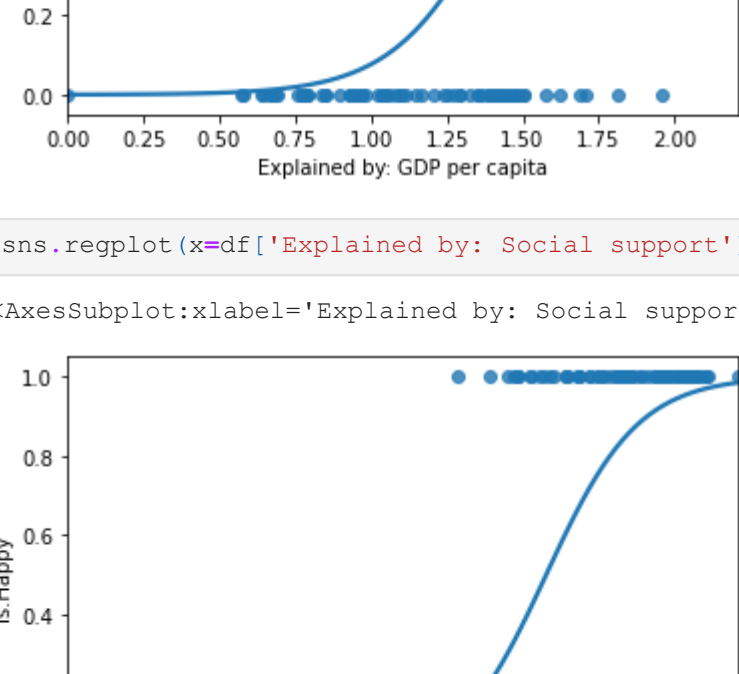
```
df['Is.Happy'].value_counts()
```

Out[65]:

```
1    74
0    72
Name: Is.Happy, dtype: int64
```


In [105.. sns.regplot(x=df['Explained by: GDP per capita'], y=y, data=df, logistic=True, ci=None)

Out[105.. <AxesSubplot:xlabel='Explained by: GDP per capita', ylabel='Is.Happy'>




In [106.. sns.regplot(x=df['Explained by: Social support'], y=y, data=df, logistic=True, ci=None)

Out[106.. <AxesSubplot:xlabel='Explained by: Social support', ylabel='Is.Happy'>



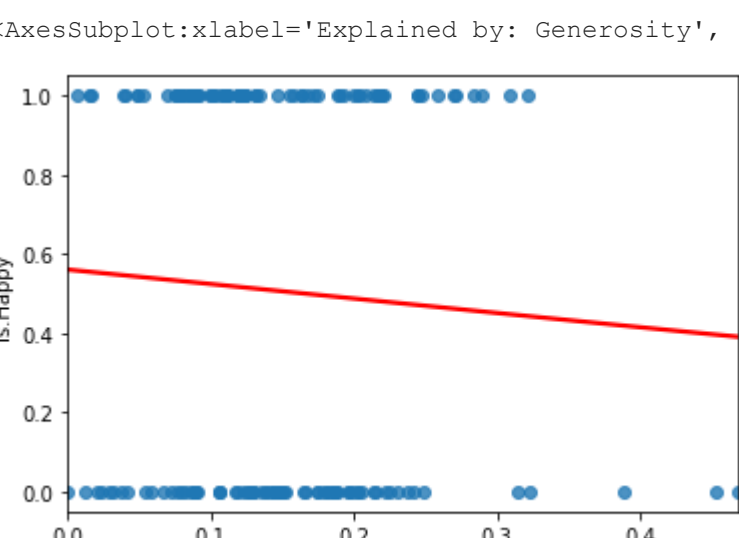
In [107.. sns.regplot(x=df['Explained by: Healthy life expectancy'], y=y, data=df, logistic=True, ci=None)

Out[107.. <AxesSubplot:xlabel='Explained by: Healthy life expectancy', ylabel='Is.Happy'>



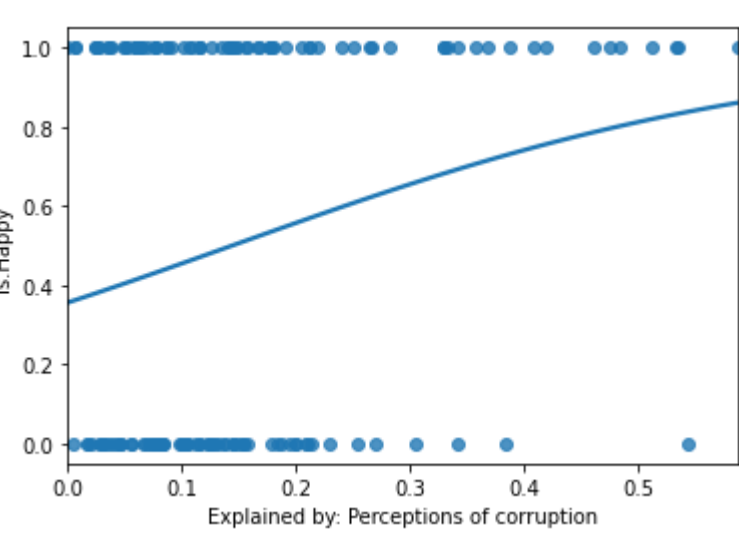
In [108.. sns.regplot(x=df['Explained by: Freedom to make life choices'], y=y, data=df, logistic=True, ci=None)

Out[108.. <AxesSubplot:xlabel='Explained by: Freedom to make life choices', ylabel='Is.Happy'>



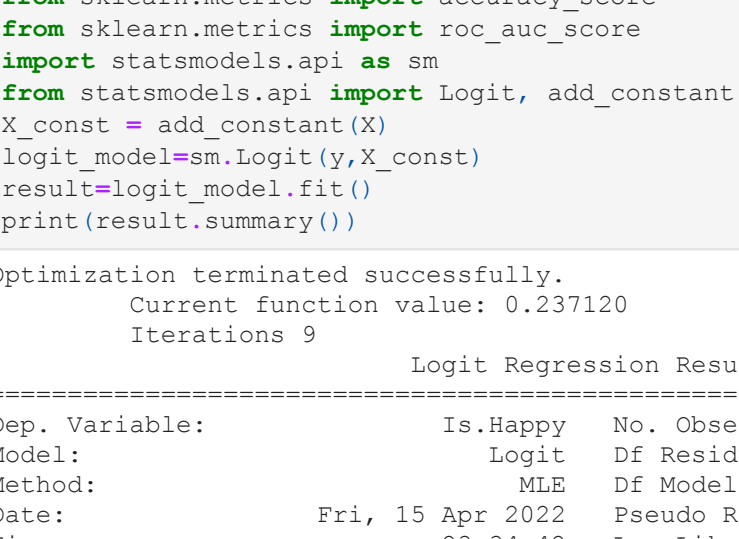
In [109.. sns.regplot(x=df['Explained by: Generosity'], y=y, data=df, line\_kws={'color': 'red'}, logistic=True, ci=None)

Out[109.. <AxesSubplot:xlabel='Explained by: Generosity', ylabel='Is.Happy'>



In [111.. sns.regplot(x=df['Explained by: Perceptions of corruption'], y=y, data=df, logistic=True, ci=None)

Out[111.. <AxesSubplot:xlabel='Explained by: Perceptions of corruption', ylabel='Is.Happy'>



In [116..

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import precision_score, recall_score, f1_score
from sklearn.datasets import make_classification
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score
import statsmodels.api as sm
logit_model=sm.Logit(y,X_const)
result=logit_model.fit()
print(result.summary())

Optimization terminated successfully.
Current function value: 0.237120
Iterations 9
```

Logit Regression Results						
Dep. Variable:	Is.Happy	No. Observations:	146			
Model:	Logit	Df Residuals:	139			
Method:	MLE	Df Model:	6			
Date:	Fri, 15 Apr 2022	Pseudo R-squ:	0.6579			
Time:	23:34:42	Log-Likelihood:	-34.620			
Converged:	True	LL-Null:	-101.19			
Covariance Type:	nonrobust	LLR p-value:	2.813e-26			
	coef	std err	z	P> z	[0.025	0.975]
const	-19.5489	3.809	-5.132	0.000	-27.015	-12.083
x1	1.1457	1.794	0.639	0.523	-2.370	4.661
x2	8.6156	2.612	3.298	0.001	3.496	13.735
x3	8.2294	3.923	2.098	0.036	0.541	15.918
x4	11.1951	3.192	3.507	0.000	4.939	17.451
x5	-5.1569	4.189	-1.231	0.218	-13.366	3.053
x6	-4.6588	3.378	-1.379	0.168	-11.280	1.963

In [117..

```
feature = ['Explained by: Social support','Explained by: Healthy life expectancy', 'Explained by: Freedom to ma
X1 = df[feature]
X1_const = add_constant(X1)
logit_model=sm.Logit(y,X1_const)
result=logit_model.fit()
print(result.summary2())

Optimization terminated successfully.
Current function value: 0.248417
Iterations 9
```

```
recall: 0.96

print('\nClassification Report:')
print(classification_report(y_test, y_pred))

import seaborn as sns
mat=confusion_matrix(y_test, y_pred)
sns.heatmap(mat,annot=True)
plt.xlabel('true label')
plt.ylabel('predicted label')
plt.title('Confusion Matrix', weight='bold')

Classification Report:
precision    recall  f1-score   support
```

In [119..

```
X_train,X_test,y_train,y_test=train_test_split(X1,y,test_size=0.3,random_state=0)
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_pred = logreg.predict(X_test)
print('Accuracy of logistic regression classifier on test set: {:.2f}'.format(logreg.score(X_test, y_test)))

Accuracy of logistic regression classifier on test set: 0.89
```

In [120..

```
from sklearn import metrics
y_pred = logreg.predict(X_test)
print(f'Accuracy Score:\n{accuracy_score(y_test, y_pred):0.3f}')
probs = logreg.predict_proba(X_test)
print('ROC AUC Score:')
print(roc_auc_score(y_test, probs[:,1]))
print("Precision:",metrics.precision_score(y_test, y_pred))
print("Recall:",metrics.recall_score(y_test, y_pred))

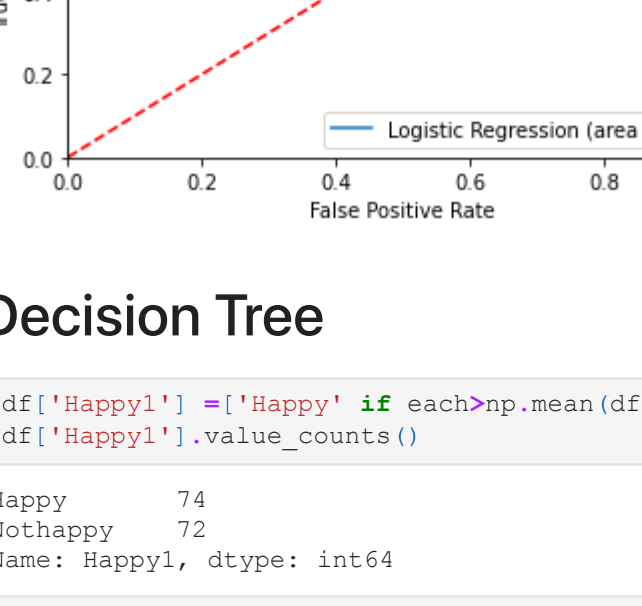
Accuracy Score:
0.886
ROC AUC Score:
0.9178947368421053
Precision: 0.8571428571428571
Recall: 0.96
```

In [121..

```
print('\nClassification Report:')
print(classification_report(y_test, y_pred))
import seaborn as sns
mat=confusion_matrix(y_test, y_pred)
sns.heatmap(mat,annot=True)
plt.xlabel('predicted label')
plt.ylabel('true label')
plt.title('Confusion Matrix', weight='bold')
```

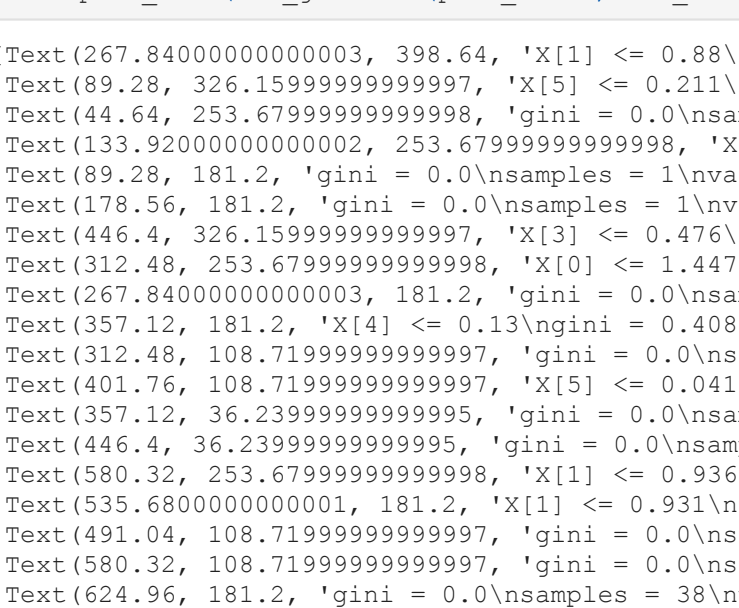
Classification Report:				
	precision	recall	f1-score	support
0	0.94	0.79	0.86	19
1	0.86	0.96	0.91	25
accuracy			0.89	44
macro avg	0.90	0.87	0.88	44
weighted avg	0.89	0.89	0.88	44

Out[121.. Text(0.5, 1.0, 'Confusion Matrix')



In [122..

```
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, logreg.predict_proba(X_test)[:,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc='lower right')
plt.savefig('Log_ROC')
plt.show()
```



# Decision Tree

In [79]: df['Happy1'] = ['Happy' if each>np.mean(df.Score) else 'Nothappy' for each in df.Score]
df['Happy1'].value\_counts()

Out[79]:

```
Happy      74
Nothappy   72
Name: Happy1, dtype: int64
```

In [90]:

```
predictors = df[['Explained by: GDP per capita', 'Explained by: Social support', 'Explained by: Healthy life ex
targets = df['Happy1']
pred_train, pred_test, tar_train, tar_test = train_test_split(predictors, targets, test_size=.4)
clf_gini = DecisionTreeClassifier(criterion='gini', random_state=0)

clf_gini.fit(pred_train,tar_train)
```

Out[90]: DecisionTreeClassifier(random\_state=0)

In [91]:

```
y_pred_gini = clf_gini.predict(pred_test)
```

In [92]:

```
plt.figure(figsize=(12,8))

from sklearn import tree

tree.plot_tree(clf_gini.fit(pred_train, tar_train))
```

Out[92]:

[Text(89.28, 326.15999999999997, 'X[5] <= 0.211\ngini = 0.062\nsamples = 31\nvalue = [1, 30]'),
Text(44.64, 253.67999999999998, 'gini = 0.0\nsamples = 12\nvalue = [0, 29]'),
Text(133.92000000000002, 253.67999999999998, 'X[0] <= 1.025\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(89.28, 181.2, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(178.56, 181.2, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(446.4, 326.15999999999997, 'X[3] <= 0.476\ngini = 0.27\nsamples = 56\nvalue = [47, 9]'),
Text(312.48, 253.67999999999998, 'X[0] <= 1.447\ngini = 0.486\nsamples = 12\nvalue = [5, 7]'),
Text(267.84000000000003, 181.2, 'gini = 0.0\nsamples = 5\nvalue = [0, 5]'),
Text(312.48, 181.2, 'X[1] <= 0.13\ngini = 0.408\nsamples = 7\nvalue = [5, 2]'),
Text(401.76, 108.71999999999997, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(491.04, 108.71999999999997, 'X[5] <= 0.041\ngini = 0.444\nsamples = 3\nvalue = [1, 2]'),
Text(357.12, 36.239999999999995, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(446.4, 36.239999999999995, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(580.32, 253.67999999999998, 'X[1] <= 0.936\ngini = 0.087\nsamples = 44\nvalue = [42, 2]'),
Text(535.68000000000001, 181.2, 'X[1] <= 0.931\ngini = 0.444\nsamples = 6\nvalue = [4, 2]'),
Text(491.04, 108.71999999999997, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(580.32, 108.71999999999997, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(624.96, 181.2, 'gini = 0.0\nsamples = 38\nvalue = [38, 0]')]

The best predictor of happiness score here would be social support